

Neural Networks Assignment 1

1. A1. DATA PREPROCESSING AND PREPARATION

Question 1. a) Data preprocessing

To properly assess the data, I used the *ydata-profiling* python library. It showed, that the data mainly consist two separate parts: the parameters about the users which is very highly correlated, and the measured values. There was a value error in the Z4 data, which was improperly exported from the csv file. The erroneous line was dropped. Also the non US compatible floating point format resulted in custom data exporting functions. Because all of the user data is highly self correlated, the repetitive data is not necessarily related, so the label type user data were all dropped. The neural network would be able to easily transform between these datasets, with for example a simple LUT like structure, they are not crucial for the training. The dropped data columns are: *age*, *user*, *gender*. Otherwise these data is not highly correlated to any other measured values, so it is not a clear issue to keep them in the training dataset.

Because the data will be used in a neural network, to make the training easier, the values are transformed to a $[0, 1]$ interval. It could be done with simple min-max scaling, but it is errorprone for outliers, which is the case for some of the truncated data noise. The data otherwise shows a clear gaussian distribution, with occasional multi gauss distributions. Only centering the data is not really necessary, most of the data is relatively centered. In the end I choose to Standardize the data, because of the general properties of the Standardization, and the general Gaussian structure of the data.

Question 2. b) cross-validation

To make my life easy, I used *Scikit Learn*'s *KFold* module, which allows to make the 5 fold CV automatically. It samples the training set based on random uniform sampling, which I considered balanced in this particular case. It was later verified by the results, that it was indeed balanced, as the variance of the results was low.

2.1. A2. CHOICE OF ARCHITECTURE

Question 3. a) Which one is preferable for training (loss)?

Cross -Entropy measures the difference between the predicted probabilities and the true probabilities, by measuring how well the predicted probability distribution of the model matches the actual probability distribution of the target classes. A lower Cross-Entropy value indicates a better model performance.

Mean Square Error, measures the average squared difference between the predicted and actual values. This shows the distance between the result and the actual values, but for a classification problem it is not really a good measure.

Accuracy shows the percentage of correct predictions from the network. It is a good measurement to show how well a network performs. It is not really a good measure for loss, because the output layer gives a probabilistic distribution over the possible classes, and accuracy would drop this information.

Question 4. b) How many neurons will you need at the output level?

For the multiclass type of classification because there are 5 classes, it needs 5 neurons to output. The embedding to this five class output is created with the *to_categorical* from library Keras.

Question 5. c) Choose an appropriate activation function for the hidden nodes?

For best adaptation for the hidden layer, and to have a good fit I used sigmoid activation, and fully connected layers.

Question 6. d) Which activation function will you use for the output level?

For classification models it is very common to use softmax activation. It is preferable because it normalizes the output of the last layer to obtain a probability distribution over the classes. The output values of the Softmax function represent the probabilities of each class, which always sum up to 1.

Question 7. Experiment with 3 different values for the number of neurons in the hidden layer

Number of neurons in the hidden layer	CE loss	MSE	Acc
THE1 = O	0.079 (+/- 0.001)	0.009 (+/- 0.000)	0.976 (+/- 0.000)
THE1 = (O+I)/2	0.116 (+/- 0.002)	0.014 (+/- 0.000)	0.966 (+/- 0.001)
THE1 = O+I	0.063 (+/- 0.005)	0.008 (+/- 0.001)	0.981 (+/- 0.002)

The problem correlates very quickly, so an early stop is clearly possible, although the author of this paper mistakenly set the epoch size for everything to 10, which is an extreme early stop.

Question 8. State your conclusions about (i) the number of hidden nodes, (ii) the choice of the cost function, and (iii) the convergence speed with respect to the training epochs.

The hidden nodes seem to have an impact on convergence speed, but in general they don't make a huge difference. The loss function is predecided to be CE and it seems to be effective. The best results were from when the number of hidden nodes were the highest, although computationally it is the most hardest task. The effects of the different hidden layer node counts can be seen on Figure ??.

8.1. A3. CHANGES IN TRAINING RATE AND MOMENTUM CONSTANT

Question 9. Optimize learning rate and momentum. It is clearly visible on Figure 1 that the lower

\eta	m	CE loss	MSE	Acc
0.001	0.2	0.794 (+/- 0.001)	0.105 (+/- 0.000)	0.737 (+/- 0.001)
0.001	0.6	0.654 (+/- 0.001)	0.088 (+/- 0.001)	0.781 (+/- 0.001)
0.05	0.6	0.118 (+/- 0.002)	0.014 (+/- 0.000)	0.966 (+/- 0.001)
0.1	0.6	0.095 (+/- 0.002)	0.011 (+/- 0.000)	0.972 (+/- 0.000)

learning rate results in a very slow learning rate, while the higher ones have a way better result. On the second row it is visible, that a high momentum results in a high initial correlation, which later slows down because of the low learning rate.

Question 10. Document theoretically why $m < 1$.

Using $m > 1$ would result in a divergent solution in this case, overshooting the correct solution.

10.1. A4. NORMALISATION

Question 11. Explain which regularization method (L1 or L2) is preferable for this problem.

The preferred method is L2 because it distributes the weight values more evenly across all the neurons, resulting in a smoother decision boundary and better generalization performance. This is easy to implement with Keras, but the results were not extremely overfitting due to the extreme early stop.

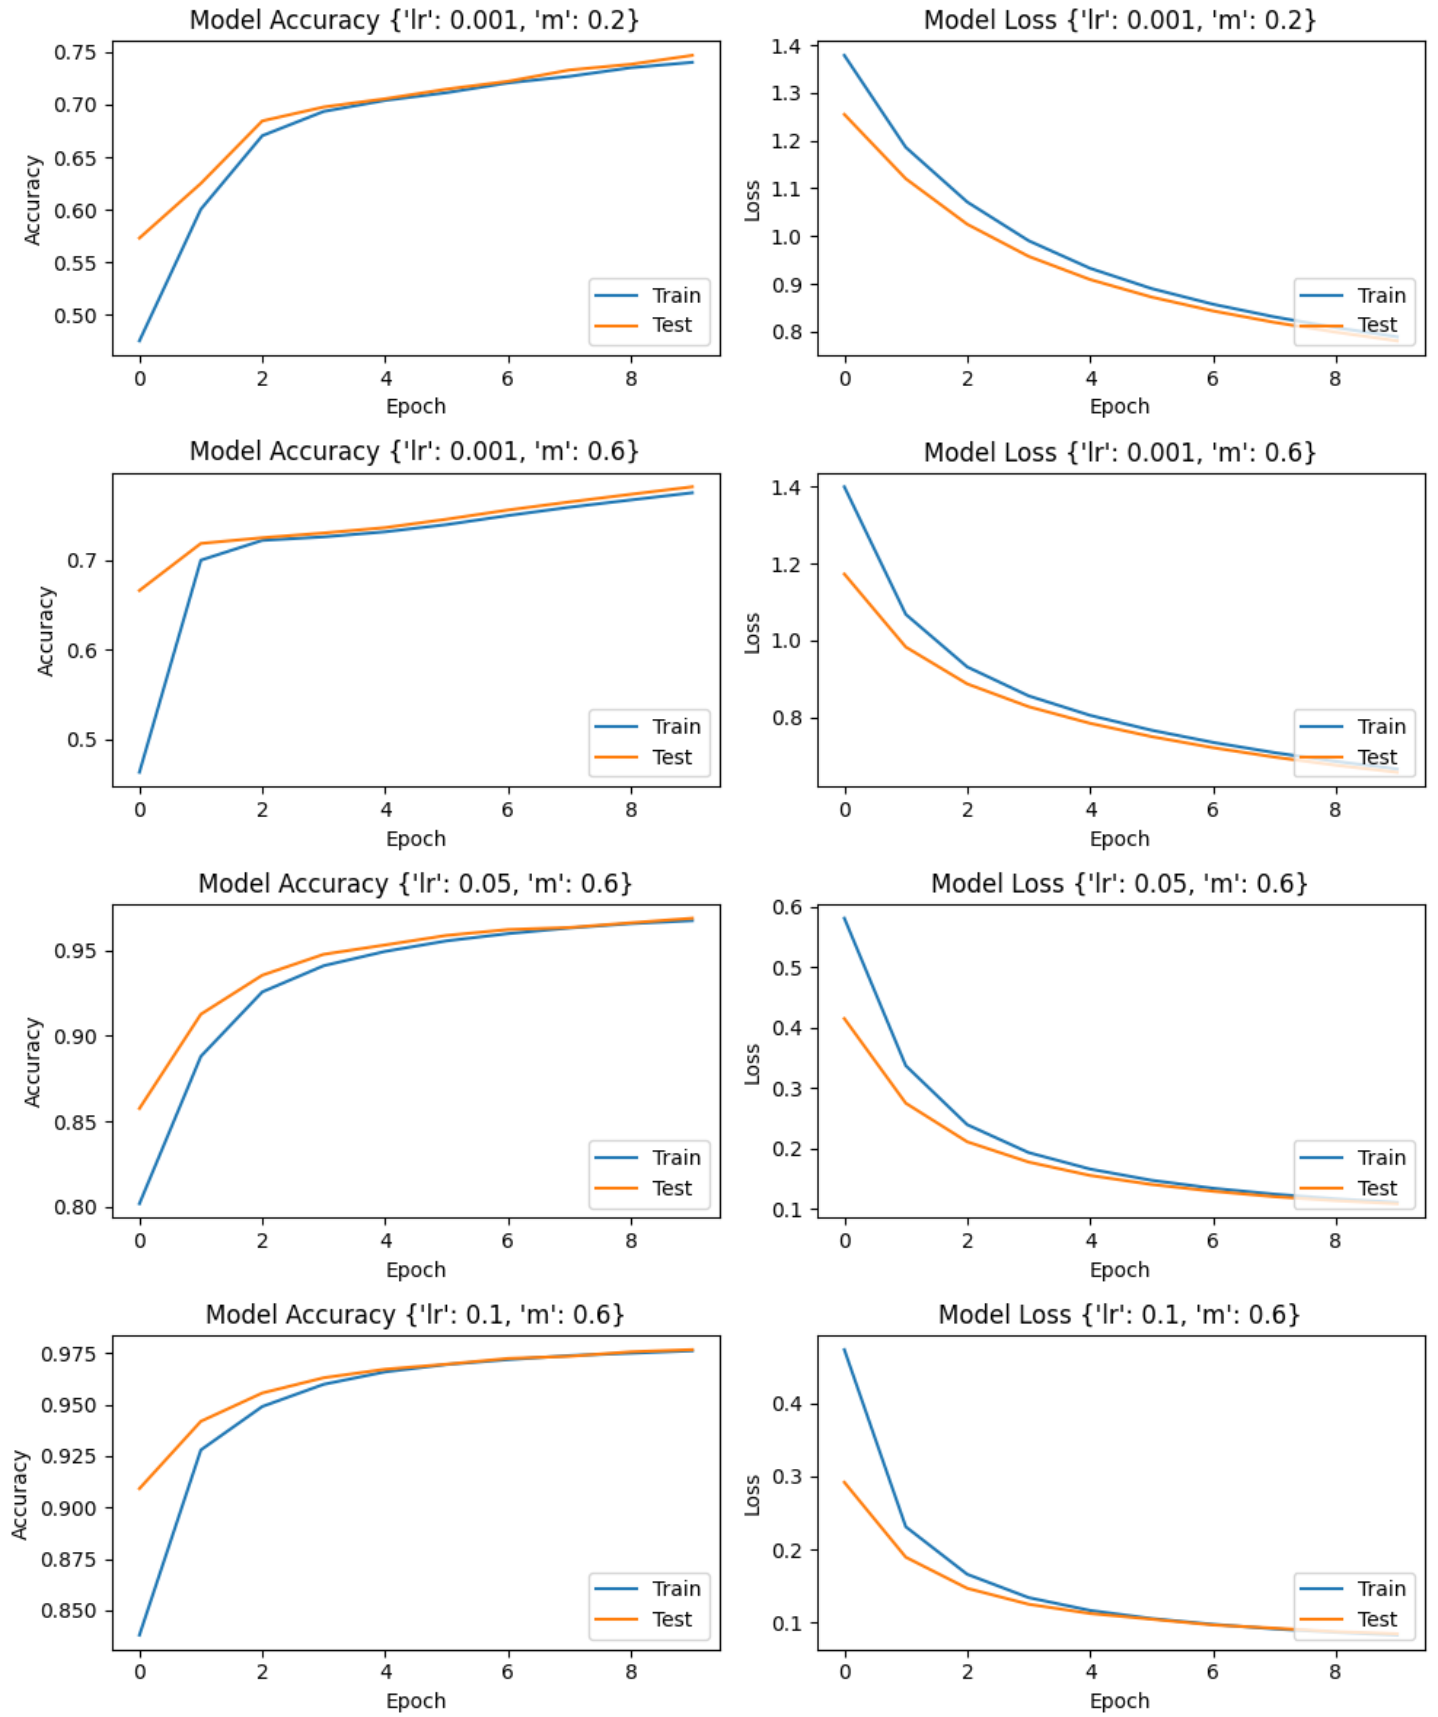


FIGURE 1. Loss and Accuracy for different number of hidden layers